

Automated Minutes of Meeting Generator

Synthetic Training Dataset — Conversation → Structured MoM

This dataset is designed for training and evaluating an automated meeting-minutes system. Each record contains a realistic meeting transcript and a reference annotation covering summary, decisions, action items, risks, and next-meeting information. Conversations intentionally include interruptions, informal language, follow-up questions, ambiguous context, and domain-specific terminology.

Dataset size: 15 meetings across software, marketing, finance, sales, research, AI, security, HR, and other scenarios.

M001 — Product Sprint Planning

Meeting type	Software Development
Participants	Rahul (PM), Priya (Frontend), Aman (Backend), Neha (QA)

Raw Meeting Transcript

Rahul: Okay, let's start with the sprint. The checkout redesign is still the priority.

Priya: The cart page is done, but the mobile layout is breaking on smaller screens.

Neha: I saw that too. Also, the coupon field throws an error when the code is expired.

Aman: That's from the API. I can return a cleaner error response instead of a generic 500.

Rahul: Good. Priya, can you finish the mobile fixes by Wednesday?

Priya: Wednesday evening should be fine.

Rahul: Aman, the API response?

Aman: I'll push it tomorrow morning.

Neha: Then I'll run regression after that.

Rahul: Let's target a Friday release if QA is clean.

Priya: One more thing—the design team changed the button colors.

Rahul: Fine, use the latest Figma file. No need for another meeting for that.

Neha: Should we include the analytics event for coupon failures?

Aman: Yes, I'll add it with the API change.

Rahul: Great. We'll review release status Thursday afternoon.

Reference Annotation

Summary: The sprint will focus on the checkout redesign. The team discussed mobile layout issues, expired coupon errors, analytics tracking, and QA before a potential Friday release.

Decisions

- Checkout redesign remains the sprint priority.
- The team will target a Friday release if regression testing passes.
- The latest Figma design will be used.

Action Items

- Priya — Fix mobile checkout layout by Wednesday evening.
- Aman — Improve expired-coupon API error response by tomorrow morning.
- Aman — Add analytics event for coupon failures.
- Neha — Run regression testing after the backend change.

Risks / Blockers

- Mobile layout issues may affect release readiness.
- Coupon error behavior currently produces a generic server error.

Next Meeting / Follow-up: Thursday afternoon release-status review

M002 — Marketing Campaign Review

Meeting type	Marketing
Participants	Neha (Marketing Lead), Arjun (Content), Simran (Social), Karan (Designer)

Raw Meeting Transcript

Neha: The October campaign is not getting enough engagement.

Simran: Reels are doing better than static posts. The last three reels got almost twice the reach.

Arjun: Then I suggest five reels a week instead of three.

Karan: I can support that, but I need scripts at least two days earlier.

Neha: Agreed. Arjun, send scripts every Monday and Thursday.

Arjun: Okay.

Simran: What about influencers? We still haven't contacted anyone.

Neha: We need two collaborations before the end of the month.

Karan: Should the influencer creatives use the same campaign identity?

Neha: Yes, but let them keep their own tone.

Simran: I'll shortlist ten creators by Friday.

Arjun: And I'll prepare the first batch of scripts tomorrow.

Neha: Good. We'll measure reach, saves, and profile visits—not just likes.

Simran: I'll add those to the weekly dashboard.

Reference Annotation

Summary: The marketing team reviewed weak October engagement and agreed to increase short-form video output, pursue influencer collaborations, and track deeper engagement metrics.

Decisions

- Increase reels from three to five per week.
- Complete two influencer collaborations during October.
- Measure reach, saves, and profile visits in addition to likes.

Action Items

- Arjun — Send scripts every Monday and Thursday.
- Arjun — Prepare the first batch of scripts tomorrow.
- Simran — Shortlist ten influencers by Friday.
- Simran — Add reach, saves, and profile visits to the weekly dashboard.

Risks / Blockers

- Higher reel frequency requires earlier script delivery.
- Influencer outreach has not started yet.

Next Meeting / Follow-up: Weekly marketing review

M003 — Database Migration Planning

Meeting type	Engineering / Infrastructure
Participants	Ananya (Engineering Manager), Dev (Backend), Karan (DevOps), Meera (QA)

Raw Meeting Transcript

Ananya: We need to finalize the production migration plan today.

Dev: The migration itself should take around three hours.

Karan: Traffic is lowest Saturday night, so that's the safest window.

Meera: What happens if the migration fails halfway?

Dev: We restore the snapshot and roll back the application schema.

Karan: I'll verify that the backup is usable before we start.

Ananya: Good. Who owns monitoring?

Dev: I do.

Ananya: And rollback execution?

Karan: I'll handle infrastructure rollback. Dev handles schema rollback.

Meera: I'll run smoke tests immediately after the migration.

Ananya: Let's freeze deployments after Friday 6 PM.

Dev: Fine.

Ananya: We need a written rollback checklist before Saturday.

Karan: I'll put that together by Friday afternoon.

Reference Annotation

Summary: The team finalized a Saturday-night production database migration with backup verification, monitoring, smoke testing, and a documented rollback procedure.

Decisions

- Migration will occur Saturday night during the lowest-traffic period.
- Deployments will be frozen after Friday 6 PM.
- Rollback responsibilities are split between infrastructure and schema changes.

Action Items

- Karan — Verify the production backup before migration.
- Dev — Monitor the migration.
- Karan — Handle infrastructure rollback if required.
- Dev — Handle schema rollback if required.
- Meera — Run post-migration smoke tests.
- Karan — Prepare the rollback checklist by Friday afternoon.

Risks / Blockers

- Migration failure could require rollback.
- The migration is estimated at three hours and may affect availability if unexpected issues occur.

Next Meeting / Follow-up: Migration execution Saturday night

M004 — Sales Pipeline Meeting

Meeting type	Sales
Participants	Rohit (Sales Lead), Pooja (Enterprise Sales), Vikram (Account Executive), Manager

Raw Meeting Transcript

Manager: Where are we on the quarterly target?

Rohit: We're at 82 percent.

Pooja: ABC Corp is the closest deal. Procurement approved the commercial terms.

Vikram: Delta Systems is waiting for their security review.

Manager: Expected closure dates?

Pooja: ABC could close this week.

Vikram: Delta probably needs another two weeks.

Manager: Fine. Prioritize ABC and don't let Delta stall.

Rohit: I'll send a pipeline update Friday.

Pooja: I need legal to review the final contract.

Manager: Send it today.

Vikram: I'll schedule a security-review call with Delta.

Manager: Good. Anything else?

Rohit: The smaller SMB pipeline is weak.

Manager: Let's discuss that next week. Focus on the enterprise deals for now.

Reference Annotation

Summary: Sales is at 82% of the quarterly target, with ABC Corp and Delta Systems as the main enterprise opportunities. The team will prioritize these deals while monitoring the weaker SMB pipeline.

Decisions

- Prioritize ABC Corp and Delta Systems.
- Defer the SMB pipeline discussion until next week.

Action Items

- Pooja — Send the final ABC contract to legal today.
- Vikram — Schedule a security-review call with Delta Systems.
- Rohit — Send a pipeline update Friday.

Risks / Blockers

- Delta Systems is blocked by a security review.
- The SMB pipeline is currently weak.

Next Meeting / Follow-up: SMB pipeline discussion next week

M005 — Customer Support Escalation

Meeting type	Customer Support
Participants	Aisha (Support Lead), Mohit (Support), Sara (Engineering), Customer Success

Raw Meeting Transcript

Aisha: We have a cluster of complaints about duplicate charges.

Mohit: I have six tickets from yesterday alone.

Sara: Are these actual duplicate charges or duplicate notifications?

Mohit: Two are definitely actual charges. The others look like notification retries.

Sara: Then I need transaction IDs for the two confirmed cases.

Aisha: I'll send those after this call.

Customer Success: One enterprise customer is threatening to pause renewal.

Aisha: That's urgent. Mohit, call them today.

Mohit: I'll do it within the next two hours.

Sara: I'll check the payment-service logs once I have the IDs.

Aisha: We also need a temporary macro for support agents so they don't promise refunds before verification.

Customer Success: I'll draft the wording.

Sara: If this is a retry bug, we should disable the affected worker.

Aisha: Only after Engineering confirms it.

Sara: Agreed.

Reference Annotation

Summary: Support escalated reports of duplicate charges. Engineering will investigate confirmed cases, while Support and Customer Success handle affected customers and interim communication.

Decisions

- Refunds should not be promised before transaction verification.
- The affected payment worker will only be disabled after Engineering confirms a retry-related defect.

Action Items

- Aisha — Send confirmed transaction IDs to Engineering.
- Mohit — Contact the enterprise customer within two hours.
- Sara — Investigate payment-service logs after receiving transaction IDs.
- Customer Success — Draft a temporary support response macro.

Risks / Blockers

- Potential duplicate-charge defect could affect renewals.
- Some reports may be notification duplicates rather than actual charges.

Next Meeting / Follow-up: Engineering investigation after transaction IDs are shared

M006 — University Project Meeting

Meeting type	Academic Project
Participants	Vibhu (Team Lead), Riya (ML), Aditya (Backend), Naman (Frontend), Faculty Mentor

Raw Meeting Transcript

Vibhu: We need the first working prototype before the next review.

Riya: The summarization model works, but the output is sometimes too long.

Faculty Mentor: Then define an evaluation metric for summary quality.

Riya: We can use ROUGE, but I also want a human evaluation score.

Aditya: The backend API is ready except authentication.

Naman: Frontend has the upload screen and transcript view.

Vibhu: What about the minutes format?

Naman: I can show summary, decisions, and action items separately.

Faculty Mentor: Add speaker names and deadlines too.

Aditya: I'll finish authentication by Friday.

Riya: I'll create the evaluation dataset by Sunday.

Naman: I'll integrate the structured response after the API is finalized.

Vibhu: Let's have everything ready for Monday's demonstration.

Faculty Mentor: Make sure the dataset contains messy conversations, not only clean examples.

Reference Annotation

Summary: The academic team reviewed an automated meeting-minutes prototype and agreed to evaluate summaries using ROUGE plus human assessment while expanding the output structure and dataset realism.

Decisions

- Evaluate summaries using ROUGE and human evaluation.
- MoM output will include summary, decisions, action items, speaker names, and deadlines.
- The dataset must include realistic, messy conversations.

Action Items

- Aditya — Complete backend authentication by Friday.
- Riya — Create the evaluation dataset by Sunday.
- Naman — Integrate structured MoM output after API finalization.
- Vibhu — Coordinate the Monday demonstration.

Risks / Blockers

- Model summaries may be excessively long.
- A clean-only dataset may not represent real meetings.

Next Meeting / Follow-up: Monday demonstration

M007 — Hiring Interview Calibration

Meeting type	Human Resources
Participants	HR Lead, Priya (Engineering), Rahul (Engineering), Recruiter

Raw Meeting Transcript

HR Lead: We need consistency across technical interviews.

Priya: Some interviewers are spending too long on the first coding question.

Rahul: I suggest a fixed 35-minute coding section.

Recruiter: What about candidates who finish early?

Rahul: Give them an extension problem, but don't change the scoring rubric.

Priya: We should also add a system-design question for senior candidates.

HR Lead: Agreed, but not for entry-level candidates.

Recruiter: I'll update the interview guide.

Priya: We need interviewer training before the new rubric goes live.

Rahul: I can run one training session next Tuesday.

HR Lead: Good. Let's pilot the rubric with five interviews first.

Recruiter: I'll select the pilot candidates.

HR Lead: We'll review the results after the pilot.

Reference Annotation

Summary: The hiring team discussed standardizing technical interviews through fixed coding time, a consistent scoring rubric, role-specific system design questions, and a pilot.

Decisions

- Use a fixed 35-minute coding section.
- Keep the scoring rubric consistent even when extension problems are used.
- Add system-design questions for senior candidates only.
- Pilot the revised rubric with five interviews.

Action Items

- Recruiter — Update the interview guide.
- Priya — Run interviewer training next Tuesday.
- Recruiter — Select five pilot candidates.

Risks / Blockers

- Inconsistent interviewer behavior may affect candidate comparisons.
- The revised rubric needs validation through the pilot.

Next Meeting / Follow-up: Pilot review after five interviews

M008 — Budget Planning Meeting

Meeting type	Finance
Participants	CFO, Finance Manager, Operations Lead, Product Lead

Raw Meeting Transcript

CFO: We need to reduce discretionary spending for the next quarter.

Finance Manager: Current discretionary spend is about 18 lakh per month.

Operations Lead: Which categories are easiest to reduce?

Finance Manager: Travel and external consulting.

Product Lead: We have two consultants supporting the migration.

CFO: Keep the migration consultant until the project reaches production.

Product Lead: That's expected in six weeks.

Operations Lead: Can we reduce travel by 25 percent?

Finance Manager: Yes, if we make client visits exception-based.

CFO: Do that.

Product Lead: We also have a planned team offsite.

CFO: Postpone it to the following quarter.

Finance Manager: I'll circulate the revised budget tomorrow.

Operations Lead: I'll update the travel policy.

Reference Annotation

Summary: Finance reviewed discretionary spending and agreed to reduce travel, postpone the team offsite, and retain migration consulting support until production.

Decisions

- Reduce travel spending by approximately 25% through exception-based approvals.
- Postpone the team offsite to the following quarter.
- Retain migration consulting support until production.

Action Items

- Finance Manager — Circulate revised budget tomorrow.
- Operations Lead — Update the travel policy.

Risks / Blockers

- Migration costs may remain elevated for six weeks.
- Travel reductions depend on effective exception handling.

Next Meeting / Follow-up: Budget review after revised forecast

M009 — Website Launch Readiness

Meeting type	Product Launch
Participants	Product Manager, Engineering Lead, QA Lead, Marketing Lead

Raw Meeting Transcript

Product Manager: Are we still comfortable launching Thursday?

QA Lead: The main checkout flow is green, but Safari has one visual issue.

Engineering Lead: It's cosmetic and doesn't block payment.

Marketing Lead: Are analytics events verified?

Engineering Lead: Purchase and signup are verified. Referral events are not.

Product Manager: Can referral tracking be fixed before Wednesday?

Engineering Lead: Yes, I'll patch it tomorrow.

QA Lead: I'll retest Safari and referral events after the patch.

Marketing Lead: We already scheduled the announcement for Thursday morning.

Product Manager: Keep it scheduled for now.

QA Lead: We need a final go/no-go Wednesday evening.

Product Manager: Agreed.

Marketing Lead: I'll prepare a backup announcement in case the launch moves.

Engineering Lead: Good.

Reference Annotation

Summary: The website remains on track for a Thursday launch, pending a Safari visual issue and referral analytics fix. A final go/no-go review is scheduled for Wednesday evening.

Decisions

- Keep the Thursday launch and announcement scheduled for now.
- Conduct a final go/no-go review Wednesday evening.

Action Items

- Engineering Lead — Fix referral analytics tomorrow.
- QA Lead — Retest Safari and referral tracking after the patch.
- Marketing Lead — Prepare a backup announcement.

Risks / Blockers

- Referral analytics is not yet verified.
- A Safari visual defect remains open.

Next Meeting / Follow-up: Wednesday evening go/no-go review

M010 — Research Literature Review

Meeting type	Research
Participants	Professor, Ananya, Dev, Rohan

Raw Meeting Transcript

Professor: We need to narrow the literature review.

Ananya: There are too many papers on generic summarization.

Rohan: We should focus on meeting summarization and dialogue structure.

Dev: I found papers using extractive and abstractive methods.

Professor: Separate them in the literature matrix.

Ananya: Should we include speech recognition papers?

Professor: Only where transcription quality directly affects summarization.

Rohan: I'll create categories for ASR, dialogue segmentation, summarization, and evaluation.

Dev: I'll collect datasets and note licensing restrictions.

Professor: Good. Don't mix benchmark results from different datasets without explaining the difference.

Ananya: I'll write the comparison section.

Professor: First draft by Friday.

Rohan: And the matrix by Wednesday.

Reference Annotation

Summary: The research team narrowed its literature review to meeting summarization and related dialogue-processing components, with separate treatment of ASR, segmentation, summarization, evaluation, datasets, and licensing.

Decisions

- Focus the review on meeting summarization and dialogue structure.
- Include speech-recognition work only when transcription quality directly affects summarization.
- Separate research categories in the literature matrix.

Action Items

- Rohan — Create literature categories and matrix by Wednesday.
- Dev — Collect datasets and document licensing restrictions.
- Ananya — Write the comparison section by Friday.

Risks / Blockers

- Benchmark results from different datasets may not be directly comparable.
- Dataset licensing may restrict project use.

Next Meeting / Follow-up: Literature matrix review Wednesday

M011 — Mobile App Incident Review

Meeting type	Incident Response
Participants	Incident Lead, Sara (Backend), Karan (SRE), QA

Raw Meeting Transcript

Incident Lead: Login failures peaked between 10:10 and 10:35.

Sara: The authentication token service started returning expired-token responses.

Karan: We saw elevated latency on the token refresh endpoint.

QA: Was there a deployment at that time?

Sara: Yes, version 4.8 went out at 10:05.

Incident Lead: Roll back?

Sara: The rollback restored login immediately.

Karan: We should still identify why refresh latency increased.

Incident Lead: Agreed. Keep the incident open.

QA: I'll reproduce the issue against the 4.8 build.

Sara: I'll inspect the token-service changes.

Karan: I'll compare latency metrics before and after deployment.

Incident Lead: Initial incident report by tomorrow noon.

Reference Annotation

Summary: A deployment caused elevated token-refresh latency and login failures. Rolling back version 4.8 restored login, but the underlying latency cause remains under investigation.

Decisions

- Keep the incident open until the token-refresh latency cause is identified.
- Use rollback as the immediate recovery mechanism for the affected release.

Action Items

- QA — Reproduce the issue against version 4.8.
- Sara — Inspect token-service changes.
- Karan — Compare latency metrics before and after deployment.
- Incident Lead — Prepare the initial incident report by tomorrow noon.

Risks / Blockers

- The root cause of token-refresh latency remains unresolved.
- A future redeployment could reproduce the incident if the cause is not identified.

Next Meeting / Follow-up: Incident review after root-cause analysis

M012 — E-commerce Recommendation Project

Meeting type	AI / E-commerce
Participants	Project Lead, ML Engineer, Backend Engineer, Data Engineer

Raw Meeting Transcript

Project Lead: The recommendation prototype is returning products, but the explanations are weak.

ML Engineer: Right now we're mostly matching keywords.

Data Engineer: We can enrich the dataset with price, rating, reviews, and product attributes.

Backend Engineer: Can we also collect current prices from multiple sites?

Project Lead: Yes, but we need to handle changing availability.

ML Engineer: Then the model should distinguish product relevance from price value.

Data Engineer: I'll design the schema for product snapshots.

Backend Engineer: I'll expose an API that returns the latest snapshot.

ML Engineer: I can build a ranking score using relevance, price, rating, and sentiment.

Project Lead: Good. Don't make the ranking purely price-based.

Data Engineer: We also need timestamps so stale prices aren't treated as current.

Project Lead: Exactly. First prototype with ten thousand product records.

Backend Engineer: I'll have the API ready Friday.

Data Engineer: Dataset schema by Wednesday.

ML Engineer: Initial ranking experiment by Monday.

Reference Annotation

Summary: The team plans to evolve the product recommendation system from keyword matching toward multi-factor ranking using relevance, price, ratings, reviews, sentiment, and freshness.

Decisions

- Separate product relevance from price value.
- Use timestamped product snapshots to handle changing information.
- Ranking will combine multiple factors rather than relying only on price.
- Initial prototype will use approximately 10,000 product records.

Action Items

- Data Engineer — Design product snapshot schema by Wednesday.
- Backend Engineer — Build latest-snapshot API by Friday.
- ML Engineer — Run initial ranking experiment by Monday.

Risks / Blockers

- Product availability and prices change over time.
- Stale snapshots could produce misleading recommendations.
- Pure keyword matching may not capture product value.

Next Meeting / Follow-up: Ranking experiment review Monday

M013 — Remote Team Retrospective

Meeting type	Team Retrospective
Participants	Team Lead, Frontend Engineer, Backend Engineer, QA Engineer, Designer

Raw Meeting Transcript

Team Lead: Let's start with what slowed us down this sprint.

Frontend: Design changes came after implementation had already started.

Designer: Fair point. I was waiting for product feedback.

Backend: I also had to wait for API requirements twice.

QA: Testing started late because the staging environment was unstable.

Team Lead: What should we change?

Designer: Freeze major designs one day before development.

Backend: And API requirements should be written before implementation.

QA: I want staging checks before every sprint demo.

Team Lead: Let's adopt all three.

Frontend: Can we also keep a shared decision log?

Team Lead: Yes. I'll create one today.

Designer: I'll add design status to the sprint board.

QA: I'll document staging checks.

Team Lead: We'll review whether this helped in the next retrospective.

Reference Annotation

Summary: The retrospective identified late design changes, incomplete API requirements, and unstable staging as major sources of delay. The team agreed on process changes and a shared decision log.

Decisions

- Freeze major designs one day before development.
- Document API requirements before implementation.
- Run staging checks before each sprint demo.
- Maintain a shared decision log.

Action Items

- Team Lead — Create shared decision log today.
- Designer — Add design status to sprint board.
- QA — Document staging checks.

Risks / Blockers

- Late requirements and unstable staging may continue causing delays if the new process is not followed.

Next Meeting / Follow-up: Next sprint retrospective

M014 — Security Review Meeting

Meeting type	Security
Participants	Security Lead, Backend Lead, DevOps, Product Manager

Raw Meeting Transcript

Security Lead: We found two medium-severity findings in the API review.

Backend Lead: One is excessive error detail in production responses.

DevOps: The other is an overly broad service-account permission.

Product Manager: Are either of these blockers for release?

Security Lead: Not if the fixes are deployed before production.

Backend Lead: I'll remove sensitive error details today.

DevOps: I'll reduce the service-account permissions tomorrow.

Security Lead: I also want logs checked after deployment.

DevOps: I can verify that.

Product Manager: Do we need another full security review?

Security Lead: No, a targeted verification is enough.

Backend Lead: I'll provide the changed API examples.

Security Lead: Good. Verification Thursday morning.

Reference Annotation

Summary: Security identified two medium-severity API issues involving error disclosure and excessive service-account permissions. Both must be fixed before production, followed by targeted verification.

Decisions

- Both security findings must be fixed before production.
- A targeted verification review is sufficient; a full review is not required.

Action Items

- Backend Lead — Remove sensitive error details today.
- DevOps — Reduce service-account permissions tomorrow.
- DevOps — Verify production logs after deployment.
- Backend Lead — Provide updated API examples.
- Security Lead — Perform targeted verification Thursday morning.

Risks / Blockers

- Production release depends on remediation of both findings.
- Incorrect permission changes could affect service functionality.

Next Meeting / Follow-up: Targeted security verification Thursday morning

M015 — Quarterly Strategy Discussion

Meeting type	Business Strategy
Participants	CEO, COO, Product Head, Sales Head, Finance Head

Raw Meeting Transcript

CEO: Revenue is growing, but the enterprise segment is taking longer to close.

Sales Head: Average enterprise cycle is now around 75 days.

Finance Head: That affects cash planning.

Product Head: We could package a smaller enterprise pilot.

COO: Would that reduce implementation time?

Product Head: Yes, because the pilot uses fewer integrations.

Sales Head: It could also lower the initial contract value.

CEO: That's acceptable if conversion improves.

Finance Head: We need a model showing pilot economics.

COO: I'll coordinate with Product and Sales on the pilot process.

Product Head: I'll define the pilot scope.

Sales Head: I'll identify five prospects for a pilot.

Finance Head: I'll model revenue and implementation costs.

CEO: Bring the proposal to next Tuesday's leadership meeting.

Reference Annotation

Summary: Leadership discussed long enterprise sales cycles and agreed to evaluate a smaller enterprise pilot model, including its economics, scope, and target prospects.

Decisions

- Evaluate a smaller enterprise pilot with fewer integrations.
- Prepare a proposal for the next Tuesday leadership meeting.

Action Items

- COO — Coordinate the pilot process with Product and Sales.
- Product Head — Define pilot scope.
- Sales Head — Identify five pilot prospects.
- Finance Head — Model pilot revenue and implementation costs.

Risks / Blockers

- Smaller pilots may reduce initial contract value.
- The proposed model may not materially shorten enterprise sales cycles.

Next Meeting / Follow-up: Tuesday leadership meeting

Recommended Training Format

For supervised fine-tuning, a practical input/output pair can be formed from each meeting. Input: the raw transcript plus an instruction such as 'Generate concise meeting minutes.' Target: the structured reference annotation. For larger experiments, create several variants of each meeting with different levels of transcript noise, speaker interruptions, transcription errors, and different output formats.

Suggested JSON fields: meeting_id, transcript, participants, summary, decisions, action_items, risks, next_meeting, minutes.